

Running Eval on a Fine-Tuned Tinker Model

Use the eval pipeline with any Tinker model path. For a fine-tuned model, this is usually the sampler weights path printed by `sample_m2_model.py`.

1. Get the Tinker Model Path

Run:

```
python sample_m2_model.py
```

Copy the printed model path. It should look like:

```
tinker://.../sampler_weights/...
```

2. Preflight the Model

Run a quick preflight before launching a full benchmark:

```
python eval-pipeline/run_benchmark_sweep.py \
--benchmark benchmarks/unified_benchmark.json \
--models "tinker://.../sampler_weights/..." \
--temperature 0.5 \
--oracle-grader \
--grader-provider tinker \
--grader-model Qwen/Qwen3-32B \
--preflight-only
```

3. Run the Full Eval

```
python eval-pipeline/run_benchmark_sweep.py \
--benchmark benchmarks/unified_benchmark.json \
--models "tinker://.../sampler_weights/..." \
--num-samples 50 \
--temperature 0.5 \
--pass-k 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 \
--workers 6 \
--samples-per-request 5 \
--delay 0 \
--oracle-grader \
--grader-provider tinker \
--grader-model Qwen/Qwen3-32B \
--run-name tinker_finetuned_<name>_t05_k1_20
```

Notes

- Replace `tinker://.../sampler_weights/...` with the actual Tinker model path.
- Keep the model path in quotes because it contains `:` and `/`.
- `--temperature 0.5` means the run samples at temperature 0.5.
- `--num-samples 50` gives 50 samples per benchmark problem.
- `--pass-k 1 ... 20` computes `pass@k` for all `k` from 1 through 20 using the same 50 samples.

- --samples-per-request 5 asks Tinker for up to 5 same-prompt completions per request when supported.